



EECS 370 - Lecture 3



LC2K

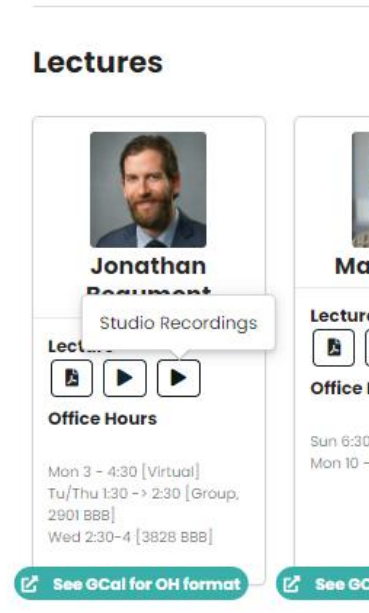


Announcements

- Lab 1 assignment due ~~Sunday~~ (extended to Thursday if you missed it)
 - No attendance required or pre-lab quiz for lab 1
 - Later labs will be due **in-person, at the end of lab**
- Pre-lab quiz for lab 2 posted on **Gradescope, due Thursday**
- P1 posted by tomorrow
 - First part due next Thursday
 - You'll have everything you need after today

Reminder: Studio Recordings

- If you're watching lectures asynchronously...
- I have studio recordings
 - Much better quality than lecture recordings



Instruction Set Architecture (ISA) Design Lectures

“People who are really serious about software should make their own hardware.” — Alan Kay

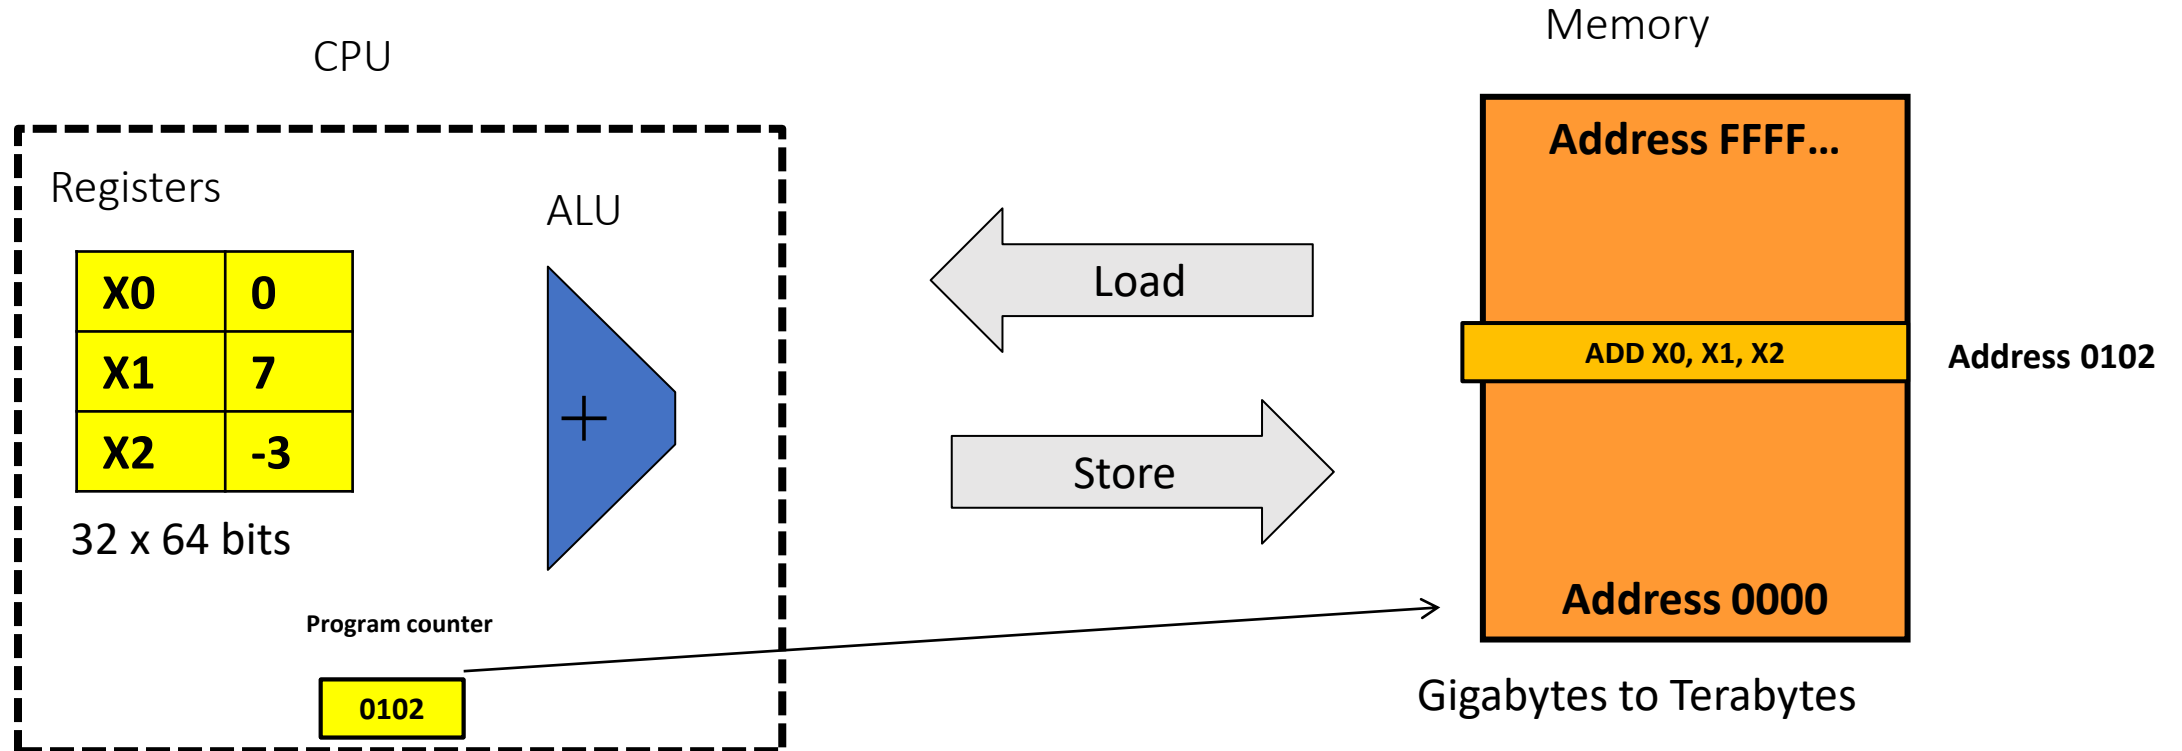
- Lecture 2: ISA - storage types, binary and addressing modes
- **Lecture 3 : LC2K**
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout



Reminder- System Organization

Let's execute this short program:

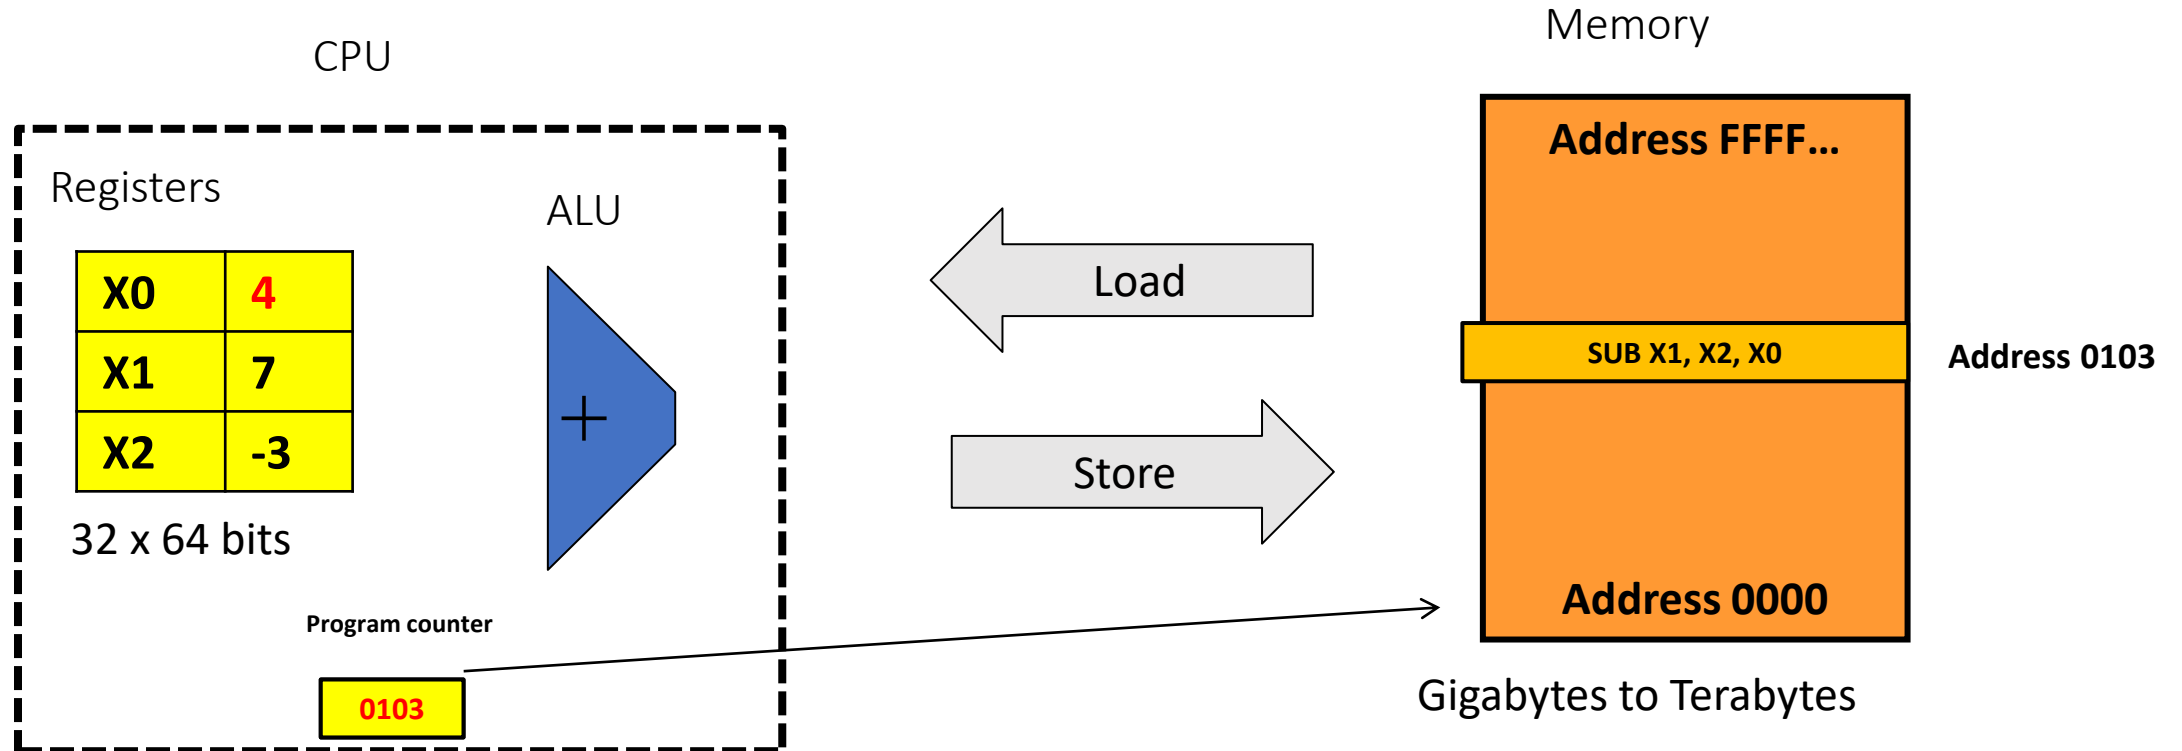
ADD X0, X1, X2
SUB X1, X2, X0



Reminder- System Organization

Let's execute this short program:


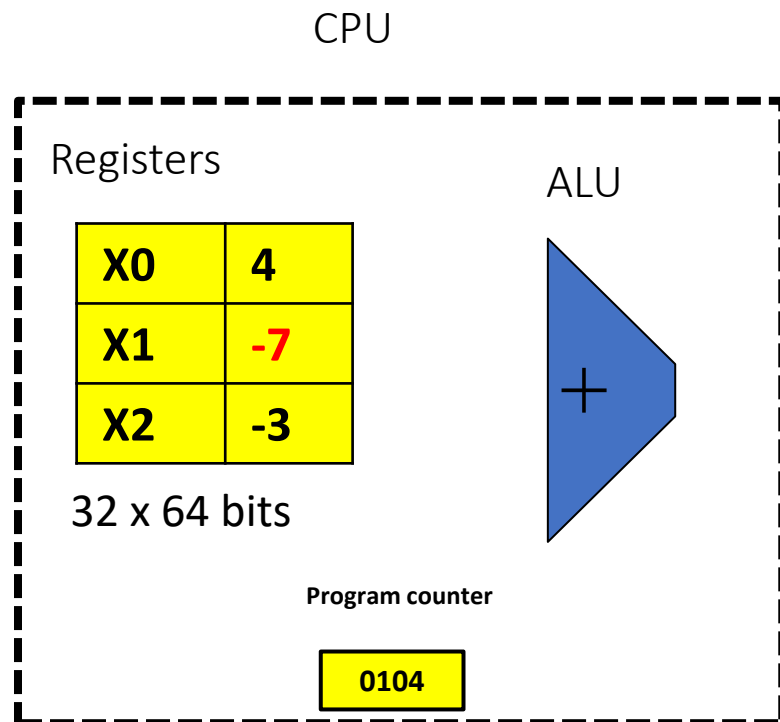
ADD X0, X1, X2
SUB X1, X2, X0



Reminder- System Organization

Let's execute this short program:

```
ADD X0, X1, X2  
SUB X1, X2, X0
```



Different Data Types

- How does memory distinguish between different data types?
 - E.g. int, int *, char, float, double
- It doesn't! It's all just 0s and 1s!
- We'll see how to encode each of these later
- Exact length depends on architectures

How is Assembly Different from C/C++?

- No data types in assembly
- Everything is 0s and 1s: up to the programmer to interpret whether these bits should be interpreted as ints, bools, chars... or even instructions themselves!

```
char c = 'a';  
c++; // c is now 'b'
```

// results in the same assembly as

```
int x = 97;  
x++; // c is now 98
```

```
x = (int) c; // this instruction has no  
effect... why?
```

Minimum Datatype Sizes

| Type | Minimum size (bits) |
|----------|---------------------|
| char | 8 |
| int | 16 |
| long int | 32 |
| float | 32 |
| double | 64 |

Representing Values in Hardware

- Unsigned integers represented as we've seen
- Chars are represented as ASCII values
 - e.g. 'a' -> 97, 'b' -> 98, '#' -> 35
- What about negative numbers?
- Fractional numbers?

Representing Negative Numbers

- There are many ways we could represent negative numbers
- Because it will eventually make our hardware simpler, the most common representation is 2's complement



Hey, Good-Looking!



2

No, not 2's *compliment*!

Two's Complement Representation

- Recall that 1101 in binary is 13 in decimal.

$$1 \ 1 \ 0 \ 1 = 8 + 4 + 1 = 13$$

$$2^3 \ 2^2 \ 2^1 \ 2^0$$

- 2's complement numbers are very similar to unsigned binary numbers.
 - The only difference is that the first number is now negative.

$$1 \ 1 \ 0 \ 1 = -8 + 4 + 1 = -3$$

$$-2^3 \ 2^2 \ 2^1 \ 2^0$$

Fun with 2's Complement Numbers

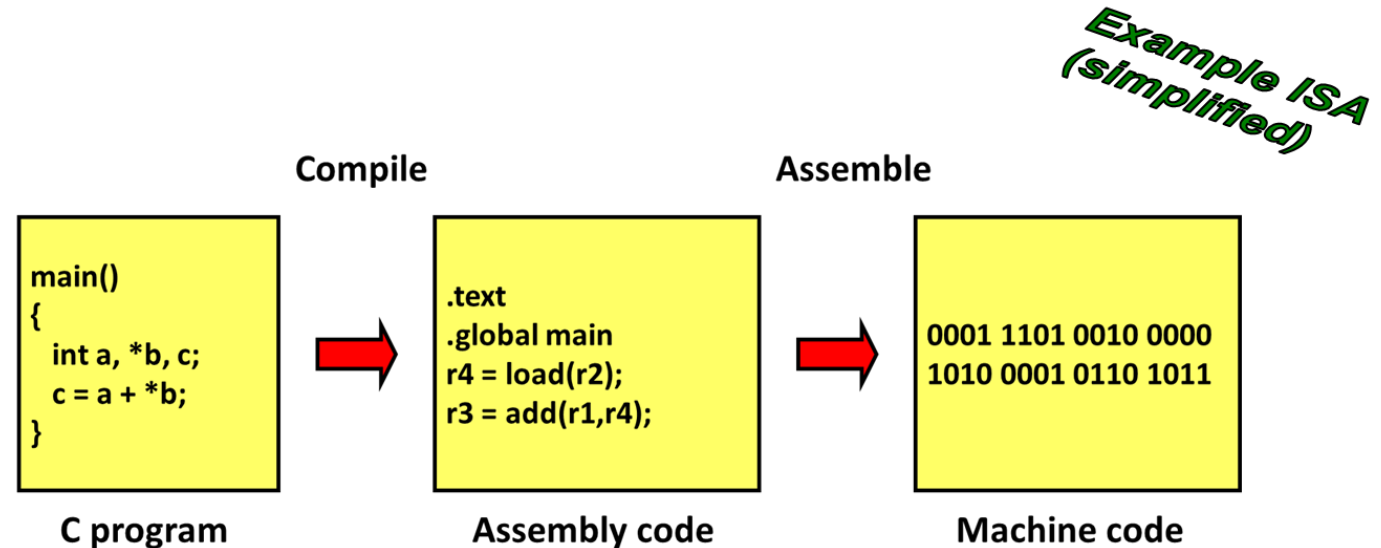
- What is the range of representation of a 4-bit 2's complement number?
 - $[-8, 7]$ (corresponding to 1000 and 0111)
- What is the range of representation of an n-bit 2's complement number?
 - $[-2^{(n-1)}, 2^{(n-1)} - 1]$
- Useful trick: You can negate a 2's complement number by inverting all the bits and adding 1.
 - 5 is represented as **0101**
 - Negate each bit: **1010**
 - Add 1: **1011** $= -8 + 2 + 1 = -5$

What about fractional numbers?

- One idea: fixed point notation
 - Have some bits represent numbers before decimal point, some bits represent numbers after decimal point
- Better idea: floating point notation
 - Inspired by scientific notation (e.g. 1.3×10^{-3})
 - Allows for larger range of numbers
 - We'll come back to this in a few lectures

Representing Instructions?

- Instructions, not just data, are stored in memory
- So, they must be expressible as numbers
- We'll look at how to encode instructions today



Agenda

- **LC2K Instruction Overview**
- Assembling LC2K into machine code
- Project 1a Overview
- Bonus Problems

LC2K Processor

- 32-bit processor
 - Instructions are 32 bits
 - Integer registers are 32 bits
- 8 registers
 - register 0 always gives the value 0
- supports 65536 words of memory (addressable space)
- 8 instructions in the following common categories:
 - Arithmetic: **add**
 - Logical: **nor**
 - Data transfer: **lw, sw**
 - Conditional branch: **beq**
 - Unconditional branch (jump) and link: **jlr**
 - Other: **halt, noop**

These are enough
instructions to express
any computation*

**(that is not limited by memory size)*

LC2K Instruction Overview: add

add 1 2 3 // r3 = r1 + r2

- Pretty self-explanatory
- What if we want to do other arithmetic operations?
 - Subtract? Same as adding, but with a negated second operand
 - Negate? In 2's complement, bitwise-NOT followed by + 1
 - Multiply? You'll figure this out for P1m

LC2K Instruction Overview: nor

`nor 1 2 3 // r3 = ~(r1 | r2)`

- Treats each source operand as binary number
- Performs bitwise NOR for each pair of bits
 - E.g. if

`r1 = 60 = 0b0000_0000_0000_0000_0000_0000_0011_1100`

`r2 = 13 = 0b0000_0000_0000_0000_0000_0000_0000_1101`

then

`r3 = 0b1111_1111_1111_1111_1111_1111_1100_0010`

- What if we want other logical operations?
 - NOT? **nor** something with itself
 - AND? Can be done using De Morgan's Law (review if needed)

LC2K Instruction Overview: lw/sw

```
// assume global variable
// is stored at address 1000
int GLOBAL;

int main() {
    GLOBAL = GLOBAL*2
}
```

```
lw    0 1 1000 // r1 = mem[1000+r0]
add   1 1 2    // r2 = 2*r1
sw    0 2 1000 // mem[1000+r0] = r2
```

- lw - "load word"
 - Loads a word (4 bytes) from a specified address into a register
- sw - "store word"
 - stores a word (4 bytes) from a register into a specified address
- Unlike add/nor, last operand here is **not** a register index
 - An **immediate** value: a number encoded directly in the instruction
- LC2K uses base+offset addressing
 - base register is first operand (if 0, then address = offset)

Non-Zero Displacement

- Consider this code:

```
struct My_Struct {  
    int tot;  
    //...  
    int val;  
};
```

```
My_Struct a;  
//...  
a.tot += a.val;
```




To load a.val...

```
lw 2 1 32  
// r1 = mem[32 + r2]  
(also need to load a.tot, add,  
then store... not shown here)
```

register file

R2 **2340**

memory



- If a register holds the starting address of "a"...
- Then the specific values needed are at a slight **offset**
- Base + Displacement**
 - reg value + immed

LC2K Instruction Overview: beq

`beq 1 2 7 // if (reg1==reg2), PC=PC+1+7`

- Remember: each line in assembly corresponds to a memory address
- "Program Counter" (PC) keeps track of address of current instruction
- Normally increments by 1
- "Branch if equal" (beq) allows us to change PC a different amount if 2 registers are equal
- Allows us to implement if/else statements, for/while loops
 - (example later)

LC2K Instruction Overview: the others

- jalr: used for function calls and returns
 - It's a bit complicated: we'll discuss later
- halt: ends the program
- noop:
 - "no operation"
 - Doesn't do anything
 - (We'll see later why this can be useful)

Note on Practical ISAs

- LC2K is made up for this class
- It's intended to be as simple as possible
 - Makes most of our projects less tedious
 - However, corresponding assembly code is **bloated**
- Practical ISAs will add many more instructions
 - Often hundreds, maybe thousands
 - Although functionally redundant, programs will be faster and easier to write


Agenda

- LC2K Instruction Overview
- **Assembling LC2K into machine code**
- Project 1a Overview
- Bonus Problems

Instruction Encoding

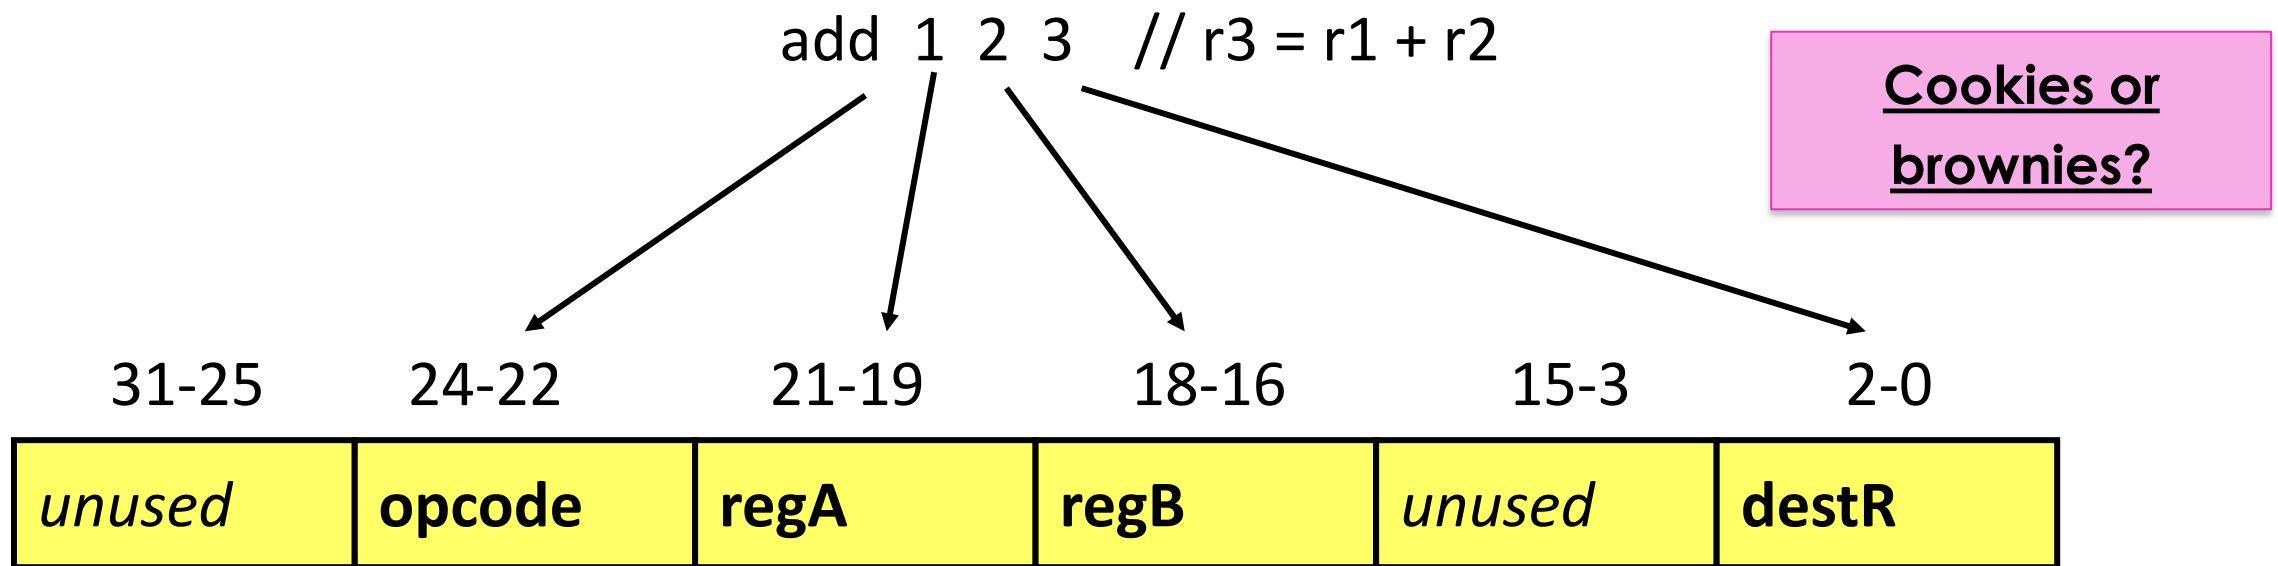
- Remember: computer doesn't understand text
 - Only understands 0s and 1s
- In order to execute our programs, assembly instructions must be converted into numbers
 - Corresponding numbers called the **machine code**
- Let's see how this is done with LC2K instructions

*Example ISA
(simplified)*



Instruction Encoding


- Instruction set architecture defines the mapping of assembly instructions to machine code




Poll: If we wanted to extend the operation code (opcode) to use all the leading unused bits, how many operations could be supported?

Instruction Formats

- Tells you which bit positions mean what
- R (register) type instructions (add, nor)




- I (immediate) type instructions (lw, sw, beq)



Instruction Formats

- J-type instructions (jalr)



- O type instructions (halt, noop)



Bit Encodings

- Most significant bits (besides unused 31-25) consist of the operation code or **opcode**
 - Indicates what type of operation
 - LC2K has 8 instructions, so we need $\log_2 8 = 3$ bits for the opcode
- Opcode encodings
 - add (000), nor (001), lw (010), sw (011), beq (100), jalr (101), halt (110), noop (111)
- Register values
 - 8 registers, so $\log_2 8 = 3$ bits for each register index
 - Just encode the register number (r2 = 010)
- Immediate values
 - Just encode the values in **2's complement format**

Next Time

- The ARM ISA

Extra Problems





Agenda

- LC2K Instruction Overview
- Assembling LC2K into machine code
- Project 1a Overview
- **Bonus Problems**



Extra Problem 1

- Compute the encoding in Hex for:
 - add 3 7 3 (r3 = r3 + r7) (add = 000)
 - sw 1 5 67 (M[r1+67] = r5) (sw = 011)



Extra Problem 1

- Compute the encoding in Hex for:
 - add 3 7 3 (r3 = r3 + r7) (add = 000)
 - sw 1 5 67 (M[r1+67] = r5) (sw = 011)

| 31-25 | 24-22 | 21-19 | 18-16 | 15-3 | 2-0 |
|---------------|---------------|-------------|-------------|---------------|--------------|
| <i>unused</i> | opcode | regA | regB | <i>unused</i> | destR |
| 0000000 | 000 | 011 | 111 | 000...000 | 011 |

| 31-25 | 24-22 | 21-19 | 18-16 | 15-0 |
|---------------|---------------|-------------|-------------|------------------|
| <i>unused</i> | opcode | regA | regB | offset |
| 0000000 | 011 | 001 | 101 | 0000000001000011 |

Extra problem 2

```
loop    lw      0    1    one
        add     1    1    1
        sw      0    1    one
        halt
one     .fill 1
```

Poll: What's the first line in binary?

- What does that program do?
- Be aware that a beq uses PC-relative addressing.
 - Be sure to carefully read the example in project 1.

