


# EECS 370 - Lecture 5

## C to Assembly



# Announcements

- P1a due Thursday
- Pre-lab 3 quiz due Thursday on Gradescope
- Reminder lab assignments must be completed in your assigned lab
  - Group assignments sent out later this week
  - Fill out admin form if you are sick / have a conflict
- HW 1 will be posted in the next day, due in two weeks

# Agenda—last lecture

- ARM overview and basic instructions
- Memory instructions
  - **Handling multiple data widths**
- Sample Problems

# Load Instruction Sizes

How much data is retrieved from memory at the given address?

Desired amount of data to transfer?	Operation	Unused bits in register?	Example
64-bits (double word or whole register)	LDUR (Load unscaled to register)	N/A	0x <b>FEDC_BA98_7654_3210</b>
16-bits (half-word) into lower bits of reg	LDURH	Set to zero	0x0000_0000_0000_ <b>3210</b>
8-bits (byte) into lower bits of reg	LDURB	Set to zero	0x0000_0000_0000_00 <b>10</b>
32-bits (word) into lower bits of reg	LDURSW (load <b>signed</b> word)	Sign extend (0 or 1 based on most significant bit of transferred word)	0x0000_0000_ <b>7654_3210</b> or 0xFFFF_FFFF_ <b>F654_3210</b> (depends on bit 31)

# But wait...

```
int my_big_number = -534159618; // 0xE0295EFE in 2's complement
```

- If I want to store this number in memory... should it be stored like this?


FE	2604
5E	2605
29	2606
E0	2607

- ... or like this?

E0	2604
29	2605
58	2606
FE	2607

# Big Endian vs. Little Endian

- Endian-ness: ordering of bytes within a word
  - Little – Bigger address holds more significant bits
  - Big – Opposite, smaller address hold more significant bits
  - The Internet is big endian, x86 is little endian, LEG and ARMv8 can switch
    - But in general assume little endian. (Figures from Wikipedia)




# Agenda—last lecture

- ARM overview and basic instructions
- Memory instructions
  - Handling multiple data widths
- **Sample Problems**

# Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STUR  X3, [X5, #100]
STURB X4, [X5, #102]
```



Memory  
(each location is 1 byte)


little endian	
0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107



# Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STUR  X3, [X5, #100]
STURB X4, [X5, #102]
```



Memory  
(each location is 1 byte)

little endian	
0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

# Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STUR  X3, [X5, #100]
STURB X4, [X5, #102]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory  
(each location is 1 byte)

little endian

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

# Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STUR  X3, [X5, #100]
STURB X4, [X5, #102]
```

register file	
X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory  
(each location is 1 byte)

little endian	
0xFF	100
0x00	101
0x00	102
0x00	103
0x00	104
0x00	105
0x00	106
0x00	107

# Example Code Sequence

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STUR  X3, [X5, #100]
STURB X4, [X5, #102]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory  
(each location is 1 byte)

little endian


0xFF	100
0x00	101
0x02	102
0x00	103
0x00	104
0x00	105
0x00	106
0x00	107

# Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STURB X3, [X5, #103]
LDURSW      X4, [X5,
#100]
```

*We shown the registers as blank. What do they actually contain before we run the snippet of code?*



Memory  
(each location is 1 byte)


little endian

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STURB X3, [X5, #103]
LDURSW      X4, [X5,
#100]
```



Memory  
(each location is 1 byte)

little endian

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STURB X3, [X5, #103]
LDURSW      X4, [X5,
#100]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory  
(each location is 1 byte)

little endian

0x02	100
0x03	101
0xFF	102
0x05	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

## Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STURB X3, [X5, #103]
LDURSW      X4, [X5,
#100]
```

register file

X3	0x0000000000000000FF
X4	0xE5FF06C205FF0302

Memory  
(each location is 1 byte)

little endian

0x02	100
0x03	101
0xFF	102
0xFF	103
0xC2	104
0x06	105
0xFF	106
0xE5	107



# Example Code Sequence (2)

What is the final state of memory once you execute the following instruction sequence? (assume X5 has the value of 0)

```
LDUR  X4, [X5, #100]
LDURB X3, [X5, #102]
STURB X3, [X5, #103]
LDURSW      X4, [X5,
#100]
```

register file

X3	0x0000000000000000FF
X4	0xFFFFFFFFFFFFFFFF0302

Memory  
(each location is 1 byte)

little endian

0x02	100
0x03	101
0xFF	102
0xFF	103
0xC2	104
0x06	105
0xFF	106
0xE5	107

# Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- **Lecture 5 : Converting C to assembly – basic blocks**
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout



# Agenda

- **Memory alignment**
  - Aligning Structs
- Control flow instructions
  - C-code examples
- Extra Problems

# Calculating Load/Store Addresses for Variables

Datatype	size (bytes)
char	1
short	2
int	4
double	8

```
short    a[100];  
char     b;  
int      c;  
double   d;  
short    e;  
struct {  
    char f;  
    int  g[1];  
    char h;  
} i;
```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

$$a = 2 \text{ bytes} * 100 = 200$$

$$b = 1 \text{ byte}$$

$$c = 4 \text{ bytes}$$

$$d = 8 \text{ bytes}$$

$$e = 2 \text{ bytes}$$

$$i = 1 + 4 + 1 = 6 \text{ bytes}$$

$$\text{total} = 221, \text{ right or wrong?}$$

# Memory layout of variables

- Compilers don't like to place variables in memory arbitrarily
- As we'll see later in the course, memory is divided into fixed-sized **chunks**
  - When we load from a particular chunk, we really read the whole chunk
  - Usually an integer number of words (32 bits)
- If we read a single char (1 byte), it doesn't matter where it's placed

0x1000	0x1001	0x1002	0x1003
'a'	'b'	'c'	'd'

ldurb [x0, 0x1002]

- Reads [0x1000-0x1003], then throws away all but 0x1002, **fine**

# Memory layout of variables

- BUT, if we read a 32-bit integer word, and that word starts at 0x1002:

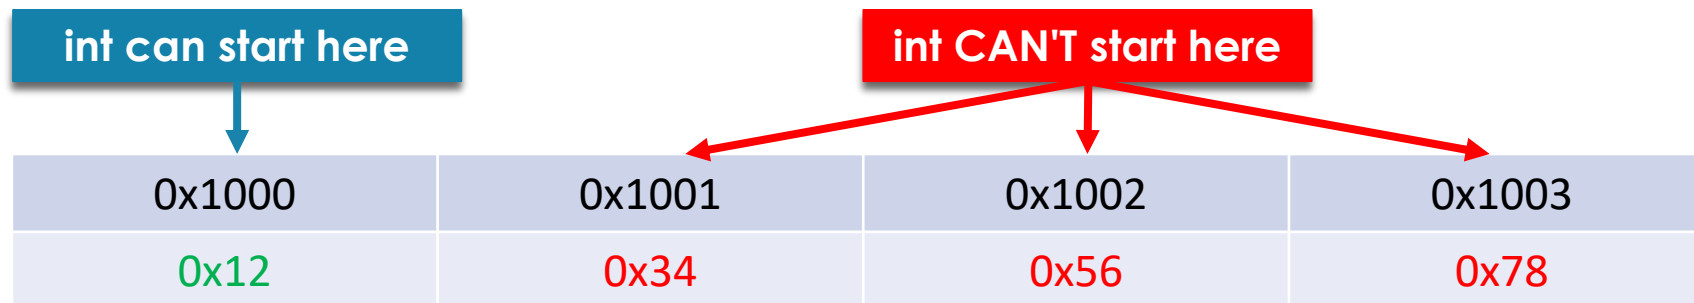
0x1000	0x1001	0x1002	0x1003
0xFF	0xFF	0x12	0x34
0x1004	0x1005	0x1006	0x1007
0x56	0x78	0xFF	0xFF

- First we need to read [0x1000-0x1003], throw away 0x1000 and 0x1001, **then** read [0x1004-0x1007]
- Need to read from memory twice! Slow! Complicated! **Bad!**


# Solution: Memory Alignment

Where can chars start?

- Most modern ISAs require that data be aligned
  - An N-byte variable must start at an address A, such that  $(A\%N) == 0$
- For example, starting address of a 32-bit **int** must be divisible by 4



- Starting address of a 16-bit **short** must be divisible by 2



# Golden Rule of Alignment

```
char  c;  
short s;  
int   i;
```

- Every (primitive) object starts at an address divisible by its size
- "Padding" is placed in between objects if needed

0x1000	0x1001	0x1002	0x1003	0x1004	0x1005	0x1006	0x1007
[c]	[padding]	[s]		[i]			

- But what about non-primitive data types?
  - Arrays? Treat as independent objects
  - Structs? Trickier...



# Agenda


- Memory alignment
  - **Aligning Structs**
- Control flow instructions
  - C-code examples
- Extra Problems

# Problem with Structs

- If we align each element of a struct according to the Golden Rule, we can still run into issues
  - E.g.: An array of structs

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

Amount of padding  
is different across  
different instances



1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	s[0].c	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]	s[1].i			

- Why is this bad?
- It makes "for" loops very difficult to write!
  - Offsets need to be different on each iteration

# Structure Alignment

- Solution: in addition to laying out each field according to Golden Rule...
  - Identify largest (primitive) field
    - Starting address of overall struct is aligned based on the largest field
    - Padded in the back so total size is a multiple of the largest primitive

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	[pad]	[pad]	[pad]	s[0].c	[pad]	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]

Guaranteed to lay  
out each instance  
identically

# Structure Example

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
}
```

**What boundary should this struct be aligned to?**

- a) 1 byte
- b) 4 bytes
- c) 12 bytes
- d) 2 bytes
- e) 19 bytes

- Assume struct starts at location 1000,
  - char w ☐ 1000
  - x[0] ☐ 1004-1007, x[1] ☐ 1008 – 1011, x[2] ☐ 1012 – 1015
  - char y ☐ 1016
  - short z ☐ 1018 – 1019

**Total size = 20 bytes!**

# Calculating Load/Store Addresses for Variables

Datatype	size (bytes)
char	1
short	2
int	4
double	8

```
short a[100];  
char b;  
int c;  
double d;  
short e;  
struct {  
    char f;  
    int g[1];  
    char h;  
} i;
```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

a = 200 bytes (100-299)

b = 1 byte (300-300)

c = 4 bytes (304-307)

d = 8 bytes (312-319)

e = 2 bytes (320-321)

struct: largest field is 4 bytes, start at 324

f = 1 byte (324-324)

g = 4 bytes (328-331)

h = 1 byte (332-332)

i = 12 bytes (324-335)

236 bytes total!! (compared to 221, originally)

# Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- Only outside structs
  - C99 forbids reordering elements inside a struct!
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- Two optimal strategies:
  - Order fields in struct by datatype size, smallest first
  - Or by largest first

# Agenda

- Memory alignment
  - Aligning Structs
- **Control flow instructions**
  - C-code examples
- Extra Problems

# ARM/LEGv8 Sequencing Instructions

- Sequencing instructions change the flow of instructions that are executed
  - This is achieved by modifying the program counter (PC)
- Unconditional branches are the most straightforward they ALWAYS change the PC and thus “jump” to another instruction out of the usual sequence
- Conditional branches

If (condition\_test) goto target\_address

*condition\_test* examines the four flags from the processor status word (SPSR)

*target\_address* is a 19 bit signed word displacement on current PC



# LEGv8 Conditional Instructions

- Two varieties of conditional branches
  1. One type compares a register to see if it is equal to zero.
  2. Another type checks the condition codes set in the status register.

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Let's look at the first type: CBZ and CBNZ
  - CBZ: Conditional Branch if Zero
  - CBNZ: Conditional Branch if Not Zero

# LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
  - The relative address is a 19 bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Example: CBNZ X3, Again
  - If X3 doesn't equal 0, then branch to label "Again"
  - "Again" is an offset from the PC of the current instruction (CBNZ)
  - Why does "25" in the above table result in PC + 100?

# LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again:      ADDI    X3, X3, #-1

-----

-----

CBNZ    X3, Again

**What is the offset?**

- a) -16
- b) -12
- c) -4
- d) -3
- e) 0

# LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again:      ADDI    X3, X3, #-1

-----

-----

CBNZ    X3, Again

- Answer = -3
  - The offset field is 19 bits signed so the bit pattern would be 111 1111 1111 1111 1101
  - Two 00's are appended to the above 19 bits and then the result would be sign-extended (with one's) to 64 bits and added to the value of PC at CBNZ
  - Why the two 00's?

# LEGv8 Conditional Instructions

- Motivation:
  - Some types of branches makes sense to check if a certain value is zero or not
    - while(a)
  - But not all:
    - if(a > b)
    - if(a == b)
  - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior

# LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations  
Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



# ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP / CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
  - N – set if the result is **negative** (i.e., bit 63 is non-zero)
  - Z – set if the result is **zero** (i.e., all 64 bits are zero)
  - ~~• C – set if last addition/subtraction had a **carry**/borrow out of bit 63~~
  - ~~• V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)~~
- Don't worry about the C and V for this class

# ARM Condition Codes Determine Direction of Branch--continued

Encoding	Name (& alias)	Meaning (integer)	Flags
0000	EQ	Equal	Z==1
0001	NE	Not equal	Z==0
0010	HS (CS)	Unsigned higher or same (Carry set)	C==1
0011	LO (CC)	Unsigned lower (Carry clear)	C==0
0100	MI	Minus (negative)	N==1
0101	PL	Plus (positive or zero)	N==0
0110	VS	Overflow set	V==1
0111	VC	Overflow clear	V==0
1000	HI	Unsigned higher	C==1 && Z==0
1001	LS	Unsigned lower or same	!(C==1 && Z==0)
1010	GE	Signed greater than or equal	N==V
1011	LT	Signed less than	N!=V
1100	GT	Signed greater than	Z==0 && N==V
1101	LE	Signed less than or equal	!(Z==0 && N==V)
1110	AL	Always	Any
1111	NV <sup>†</sup>		

Need to know the 7 with the red arrows

```
CMP X1, X2
B.LE Label1
```

For this example, we branch if X1 is >= to X2



# Conditional Branches: How to use

- CMP instruction lets you compare two registers.
  - Could also use SUBS etc.
    - That could save you an instruction.
- B.cond lets you branch based on that comparison.

- Example:

```
CMP    X1, X2  
B.GT   Label1
```

- Branches to Label1 if X1 is greater than X2.

# Agenda

- Memory alignment
  - Aligning Structs
- Control flow instructions
  - **C-code examples**
- Extra Problems

# Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

# Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

## Using Labels

```
CMP    X1, X2  
B.NE   L1  
ADD    X1, X1, #1  
B      L2  
L1: ADD    X2, X2, #1  
L2: ...
```

Note that conditions in assembly are often the inverse of the "if" condition. Why?

## Without Labels

```
CMP    X1, X2  
B.NE   3  
ADD    X1, X1, #1  
B      2  
ADD    X2, X2, #1
```

Assemblers must deal with labels and assign displacements

# Loop—Example

// assume all variables are long long integers (64 bits or 8 bytes)  
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

# of branch instructions  
=  $3 \cdot 10 + 1 = 31$

a.k.a. while-do template

```
        MOV    X1, XZR  
        MOV    X2, XZR  
Loop1:  CMPI    X1, #10  
        B.EQ   endLoop  
        LSL    X6, X1, #3  
        LDUR   X5, [X6, #100]  
        CMPI   X5, #0  
        B.LT   endif  
        ADD    X2, X2, X5  
endif:  ADDI    X1, X1, #1  
        B      Loop1  
endLoop:
```

# Agenda

- Memory alignment
  - Aligning Structs
- Control flow instructions
  - C-code examples
- **Extra Problems**

# Extra Example: Do-while Loop

// assume all variables are long long integers (64 bits or 8 bytes)  
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

# of branch instructions  
=  $2 * 10 = 20$

a.k.a. do-while template

```
                MOV    X1, XZR  
                MOV    X2, XZR  
Loop1:          LSL     X6, X1, #3  
                LDUR    X5, [X6, #100]  
                CMPI    X5, #0  
                B.LT    endif  
                ADD     X2, X2, X5  
endif:          ADDI    X1, X1, #1  
                CMPI    X1, #10  
                B.LT    Loop1  
endLoop:
```

# Extra Example: Do-while Loop

// assume all variables are long long integers (64 bits or 8 bytes)  
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

# of branch instructions  
=  $2 \times 10 = 20$

a.k.a. do-while template



# Extra Problem – For Your Reference

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in X1  
// struct {int val; struct node *next;} node;  
// struct node *ptr;
```

```
if ((ptr != NULL) && (ptr->val > 0))  
    ptr->val++;
```

# Extra Problem

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in X1
// struct {int val; struct node *next;} node;
// struct node *ptr;
```

```
if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

```
cmp r1, #0
beq Endif
ldursw r2, [r1, #0]
cmp r2, #0
b.le Endif
add r2, r2, #1
str r2, [r1, #0]
Endif : ....
```

# Extra Class Problem

- How much memory is required for the following data, assuming that the data starts at address 200 and is a 32 bit address space?

```
int a;  
struct {double b, char c, int d} e;  
char* f;  
short g[20];
```

**How much memory?**

- a)  $x < 40$  bytes
- b)  $40 < x < 50$  bytes
- c)  $50 < x < 60$  bytes
- d)  $60 < x$  bytes

# Next Time

- More C-to-Assembly
  - Function calls