

EECS 370 - Lecture 6

Function Calls



Announcements

- Project 1a due tonight
- Project 1s+m due next Thursday
- HW 1 Posted
 - Due Monday 9/22
- Let us know about exam conflicts in the **next week**
 - Form on Ed

LEGv8 Conditional Instructions

- Two varieties of conditional branches
 1. One type compares a register to see if it is equal to zero.
 2. Another type checks the condition codes set in the status register.

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Let's look at the first type: CBZ and CBNZ
 - CBZ: Conditional Branch if Zero
 - CBNZ: Conditional Branch if Not Zero

LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
 - The relative address is a 19 bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Example: CBNZ X3, Again
 - If X3 doesn't equal 0, then branch to label "Again"
 - "Again" is an offset from the PC of the current instruction (CBNZ)
 - Why does "25" in the above table result in PC + 100?

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again: ADDI X3, X3, #-1

CBNZ X3, Again

What's the offset?

- a) -16
- b) -12
- c) -4
- d) -3
- e) 0

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

Again: ADDI X3, X3, #-1

CBNZ X3, Again

- Answer = -3
 - The offset field is 19 bits signed so the bit pattern would be 111 1111 1111 1111 1101
 - Two 00's are appended to the above 19 bits and then the result would be sign-extended (with one's) to 64 bits and added to the value of PC at CBNZ
 - Why the two 00's?

LEGv8 Conditional Instructions

- Motivation:
 - Some types of branches makes sense to check if a certain value is zero or not
 - while(a)
 - But not all:
 - if(a > b)
 - if(a == b)
 - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior

LEGv8 Conditional Instructions Using FLAGS

- FLAGS: NZVC record the results of (arithmetic) operations
Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP /CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
 - N – set if the result is **negative** (i.e., bit 63 is non-zero)
 - Z – set if the result is **zero** (i.e., all 64 bits are zero)
 - ~~• C – set if last addition/subtraction had a **carry**/borrow out of bit 63~~
 - ~~• V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)~~
- Don't worry about the C and V for this class

ARM Condition Codes Determine Direction of Branch--continued

	Encoding	Name (& alias)	Meaning (integer)	Flags
→	0000	EQ	Equal	Z==1
→	0001	NE	Not equal	Z==0
	0010	HS (CS)	Unsigned higher or same (Carry set)	C==1
	0011	LO (CC)	Unsigned lower (Carry clear)	C==0
	0100	MI	Minus (negative)	N==1
	0101	PL	Plus (positive or zero)	N==0
	0110	VS	Overflow set	V==1
	0111	VC	Overflow clear	V==0
	1000	HI	Unsigned higher	C==1 && Z==0
	1001	LS	Unsigned lower or same	!(C==1 && Z==0)
→	1010	GE	Signed greater than or equal	N==V
→	1011	LT	Signed less than	N!=V
→	1100	GT	Signed greater than	Z==0 && N==V
→	1101	LE	Signed less than or equal	!(Z==0 && N==V)
→	1110	AL	Always	Any
	1111	NV [†]		

Need to know
the 7 with the
red arrows

```
CMP X1, X2
B.LE Label1
```

For this example,
we branch if X1 is
≤ to X2

Conditional Branches: How to use

- CMP instruction lets you compare two registers.
 - Could also use SUBS etc.
 - That could save you an instruction.
- B.cond lets you branch based on that comparison.

- Example:

```
CMP    X1, X2  
B.GT   Label1
```

- Branches to Label1 if X1 is greater than X2.

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - **C-code examples**
- Extra Problems

Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Using Labels

```
CMP X1, X2  
B.NE L1  
ADD X1, X1, #1  
B L2  
L1: ADD X2, X2, #1  
L2: ...
```

Note that conditions in assembly are often the inverse of the "if" condition. Why?

Without Labels

```
CMP X1, X2  
B.NE 3  
ADD X1, X1, #1  
B 2  
ADD X2, X2, #1
```

Assemblers must deal with labels and assign displacements

Loop—Example

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

of branch instructions
= $3 \cdot 10 + 1 = 31$

a.k.a. while-do template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	CMPI	X1, #10
	B.EQ	endLoop
	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	B	Loop1
endLoop:		

Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- **Lecture 6 : Converting C to assembly – functions**
- Lecture 7 : Translation software; libraries, memory layout



Agenda

- **Branching far away**
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

Branching far away

- Underlying philosophy of ISA design: **make the common case fast**
- Most branches target nearby instructions
 - Displacement of 19 bits is usually enough
- BUT what if we need to branch really far away (more than 2^{19} words)?

CBZ X15, FarLabel

- The assembler is smart enough to replace that with

CBNZ X15, L1

B FarLabel

L1:

- The simple branch instruction (B) has a 26 bit offset which spans about 64 million instructions!
- In LC2K, we can do a similar thing by using JALR instead of BEQ

Agenda

- Branching far away
- **Function calls and the call stack**
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

Implementing Functions

Poll: What's wrong
with this?

- Does this assembly code do what we need?

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main() {  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL]  
ADD     X2, X1, #1           // Inc GLOBAL  
STURW   X2, [XZR, X]        // Pass arg  
B       MULT_2              // Execute func  
  
RETURN:  
LDURSW X3, [XZR, TEMP]      // load result  
STURW   X3, [XZR, STRING]   // Pass arg  
B       PRINTF              // Execute func  
  
...  
MULT_2:  
LDURSW X1, [XZR, X]         // load arg  
ADD     X2, X1, X1          // mult by 2  
STURW   X2, [XZR, TEMP]     // return result  
B       RETURN              // return
```

Problem 1: Returning from Functions

- Branches so far have hard-coded destination

```
B.NE L1
ADD X1, X1, #1
B L2
L1: ADD X2, X2, #1
L2: ...
```

```
B.NE 3
ADD X1, X1, #1
B 2
ADD X2, X2, #1
```

- This is fine for if-statements, for-loops etc
- But functions can be called from multiple places
 - Meaning we'll return to different spots on each func call! Can't hardcode offset!

```
int func(int x) {
    printf(x * 10);
    return;
}
int helper() {
    func(7);
}
int main() {
    helper();
    func(13);
}
```

Should this return to
"helper" or "main"?

Solution: Indirect Jumps

- Indirect branches or "jumps" don't hardcode destination
- They index a register whose value holds destination

Unconditional branch	branch	B	2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register	BR	X30	go to X30	For switch, procedure return
	branch with link	BL	2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

- Use "BL" to **call a function**
 - Destination is hardcoded
 - PC +4 (return address) stored in X30
- Use "BR" to **return from a function**
 - X30 is read for return address
 - Allows us to return to different places

Solution: Indirect Jumps

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}
```

```
int GLOBAL = 6;
```

```
int main() {  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

Also don't
need "return"
labels

Now MULT_2
can return to
whatever
function
called it

```
LDURSW X1, [XZR, GLOBAL]  
ADD     X2, X1, #1           // Inc GLOBAL  
STURW   X2, [XZR, X]         // Pass arg  
BL      MULT_2               // Execute func  
  
RETURN:  
LDURSW X3, [XZR, TEMP]      // load result  
STURW   X3, [XZR, STRING]   // Pass arg  
BL      PRINTF              // Execute func  
  
...  
MULT_2:  
LDURSW X1, [XZR, X]         // load arg  
ADD     X2, X1, X1          // mult by 2  
STURW   X2, [XZR, TEMP]     // return result  
BR      X30                 // return
```

Problem 2: Passing Parameters

For any recursive functions, global variables will be overwritten

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main() {  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL]  
ADD     X2, X1, #1           // Inc GLOBAL  
STURW   X2, [XZR, X]         // Pass arg  
BL      MULT_2               // Execute func  
LDURSW  X3, [XZR, TEMP]      // load result  
STURW   X3, [XZR, STRING]    // Pass arg  
BL      PRINTF               // Execute func  
...  
MULT_2:  
LDURSW  X1, [XZR, X]         // load arg  
ADD     X2, X1, X1           // mult by 2  
STURW   X2, [XZR, TEMP]      // return result  
BR                               // return
```


Passing parameters

- Where should you put all of the parameters?
 - Registers?
 - Fast access but few in number and wrong size for some objects
 - Memory?
 - Good general solution but slow
- ARMv8 solution—and the usual answer:
 - Both
 - Put the first few parameters in registers (if they fit) (X0 – X7)
 - Put the rest in memory on the call stack— **important concept**

Call stack

- ARM conventions (and most other processors) allocate a region of memory for the “call” stack
 - This memory is used to manage all the storage requirements to simulate function call semantics
 - Parameters (that were not passed through registers)
 - Local variables
 - Temporary storage (when you run out of registers and need somewhere to save a value)
 - Return address
 - Etc.
- Sections of memory on the call stack [**stack frames**] are allocated when you make a function call, and de-allocated when you return from a function


The stack grows as functions are called

FUNCTION CALLS

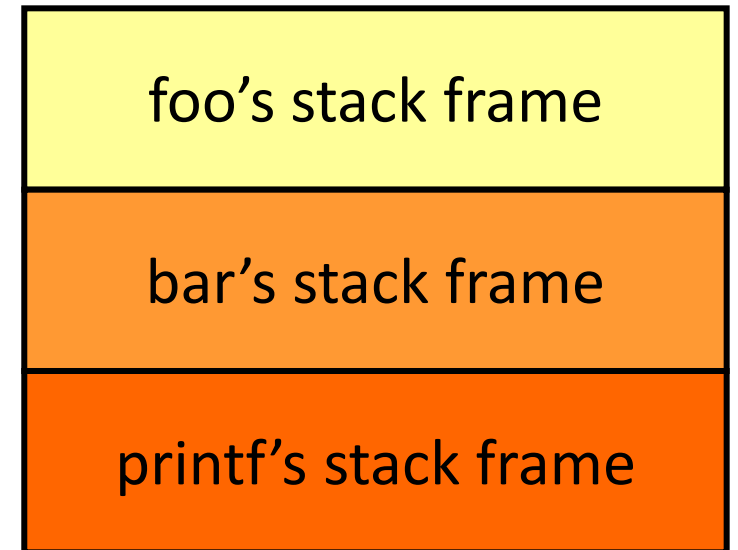
```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

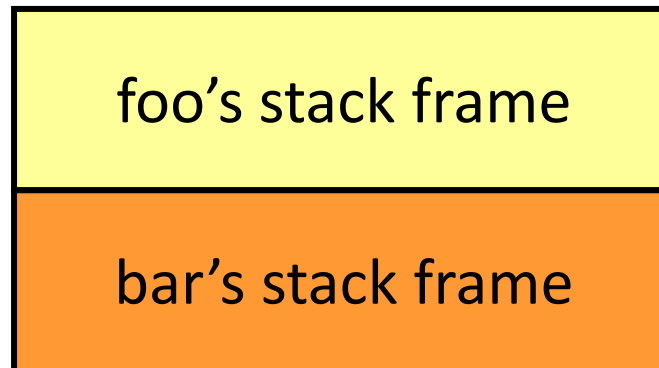
inside foo



bar calls printf



foo calls bar



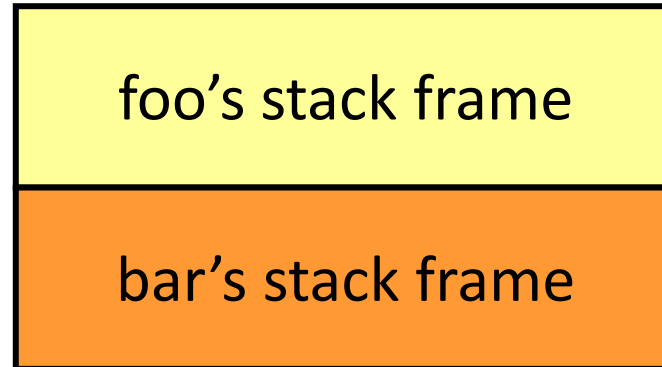
The stack shrinks as functions return

FUNCTION CALLS


```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

printf returns



bar returns



Stack frame contents

FUNCTION CALLS

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}  
  
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

foo's stack frame

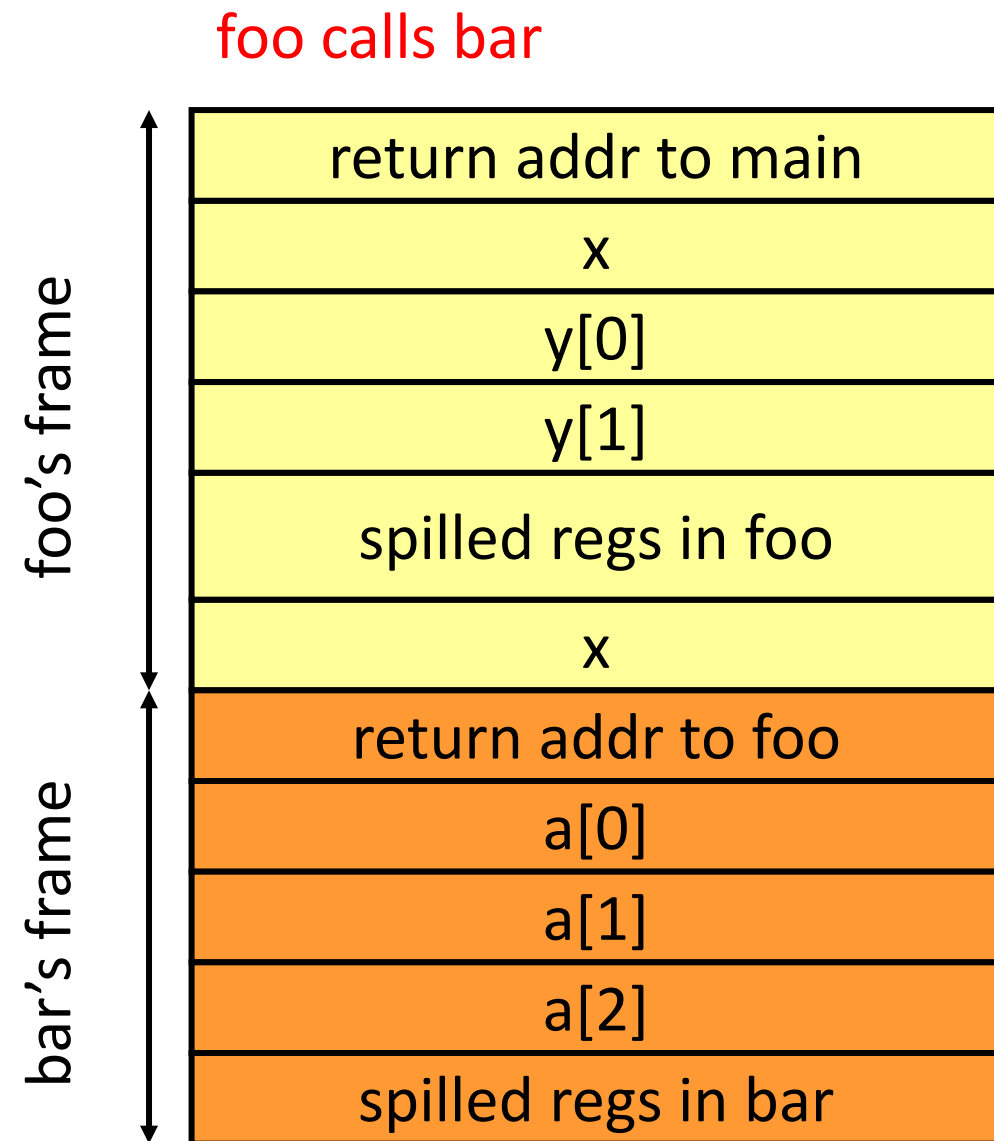
return addr to main
x
y[0]
y[1]
spilled registers in foo

Stack frame contents (2)

```
void foo()
{
    int x, y[2];
    bar(x);
}

void bar(int x)
{
    int a[3];
    printf();
}
```

Spill data—not enough room in x0-x7 for
params and also caller and callee saves



Agenda


- Branching far away
- Function calls and the call stack
- **Assigning variables to memory locations**
- Saving registers
- Caller/callee example

Review: Where do the variables go?

Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```



Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```

w goes in static, as it's a global

x goes on the stack, as it's a parameter


y goes in static, 1 copy of this!!

p goes on the stack

allocate 10 bytes on heap, ptr
set to the address

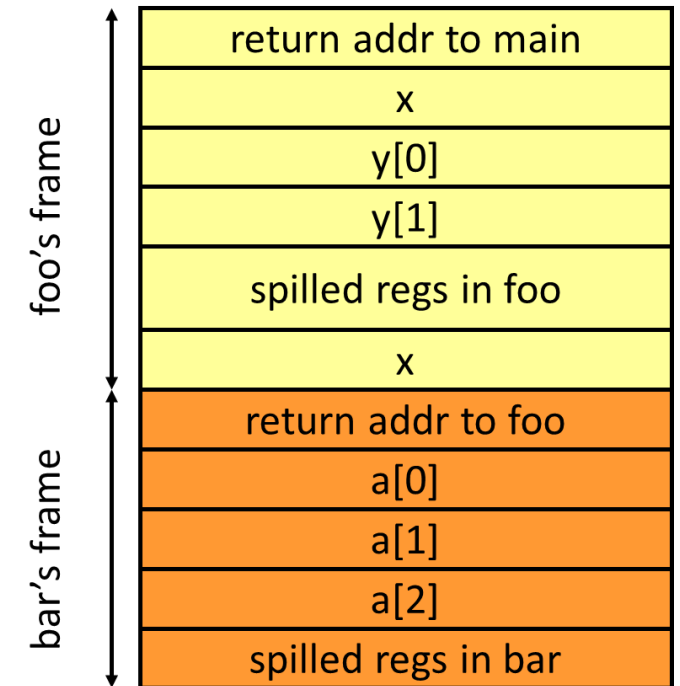
string literal "%s\n" goes in static,
implicit pointer to string on stack, p goes
on stack

The addresses of local variables
will be different depending on
where we are in the call stack



Accessing Local Variables

- Stack pointer (SP):
 - register that keeps track of current top of stack
- Compiler (or assembly writer) knows relative offsets of objects in stack
- Can access using lw/sw offsets

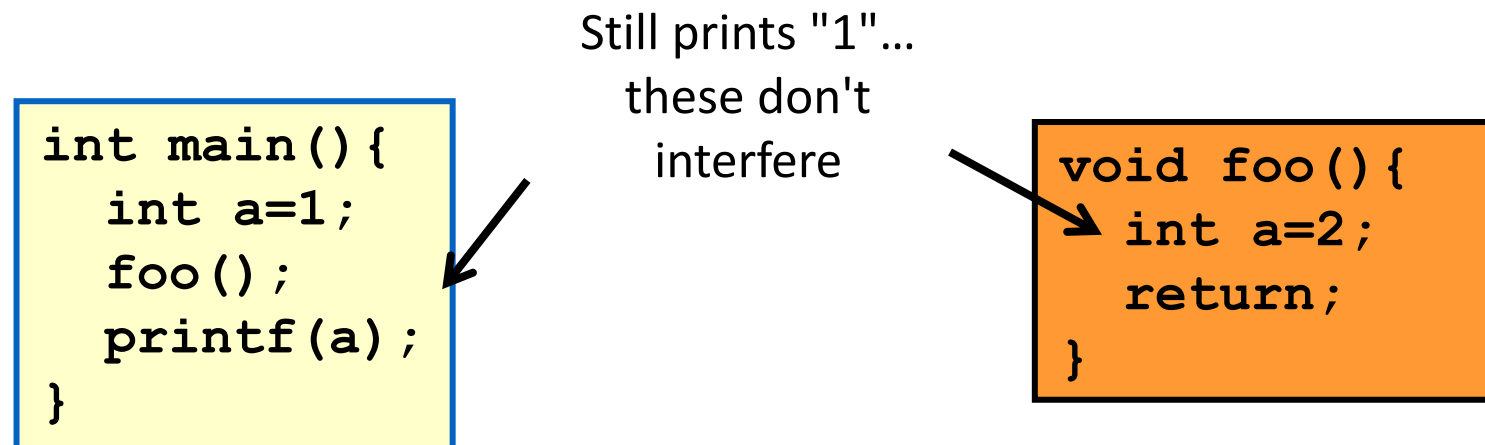


Agenda

- Branching far away
- Function calls and the call stack
- Assigning variables to memory locations
- **Saving registers**
- Caller/callee example

Problem 3: Reusing registers

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
 - Even if different function have local variables with the same name, they are independent and guaranteed not to interfere with each other!



What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
 - Called functions will overwrite registers needed by calling functions

```
main: movz X0, #1  
      bl  foo  
      bl  printf
```

foo() overwrites
X0 if we don't
do something!!

```
foo: movz X0, #2  
     br  X30
```

- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee** saved)...

```
main: movz X0, #1
      bl  foo
      bl  printf
```

```
foo:  stur X0, [stack]
      movz X0, #2
      ldur X0, [stack]
      br  X30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller** saved)...

```
main: movz X0, #1
      stur X0, [stack]
      bl  foo
      ldur X0, [stack]
      bl  printf
```

```
foo:  movz X0, #2
      br  X30
```

Another example

Original C Code

```
void foo() {
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;

    bar();

    d = a+d;
    return();
}
```

No need to
save r2/r3.
Why?

Additions for Caller-save

```
void foo() {
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;
    save r1 to stack
    save r4 to stack
    bar();
    restore r4
    restore r1
    d = a+d;
    return();
}
```

Assume bar() will
overwrite registers
holding a,d

Additions for Callee-save

```
void foo() {
    int a,b,c,d;
    save r1
    save r2
    save r3
    save r4
    a = 5; b = 6;
    c = a+1; d=c-1;
    bar();
    d = a+d;
    restore r4
    restore r3
    restore r2
    restore r1
    return();
}
```

bar() will save a,b, but
now foo() must save
main's variables

“caller-save” vs. “callee-save”

- Caller-save

- What if bar() doesn't use r1/r4?
- No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r1  
    restore r4  
    d = a+d;  
    return();  
}
```

- Callee-save

- What if main() doesn't use r1-r4?
- No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r1  
    restore r2  
    restore r3  
    restore r4  
    return();  
}
```


Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- **Caller-saved**
 - Only needs saving if value is “live” across function call
 - **Live** = contains a useful value: Assign value before function call, use that value after the function call
 - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT
live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
 - Only needs saving at beginning of function and restoring at end of function
 - Only save/restore it if function overwrites the register

Only use r1-
r4

No need to
save other
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

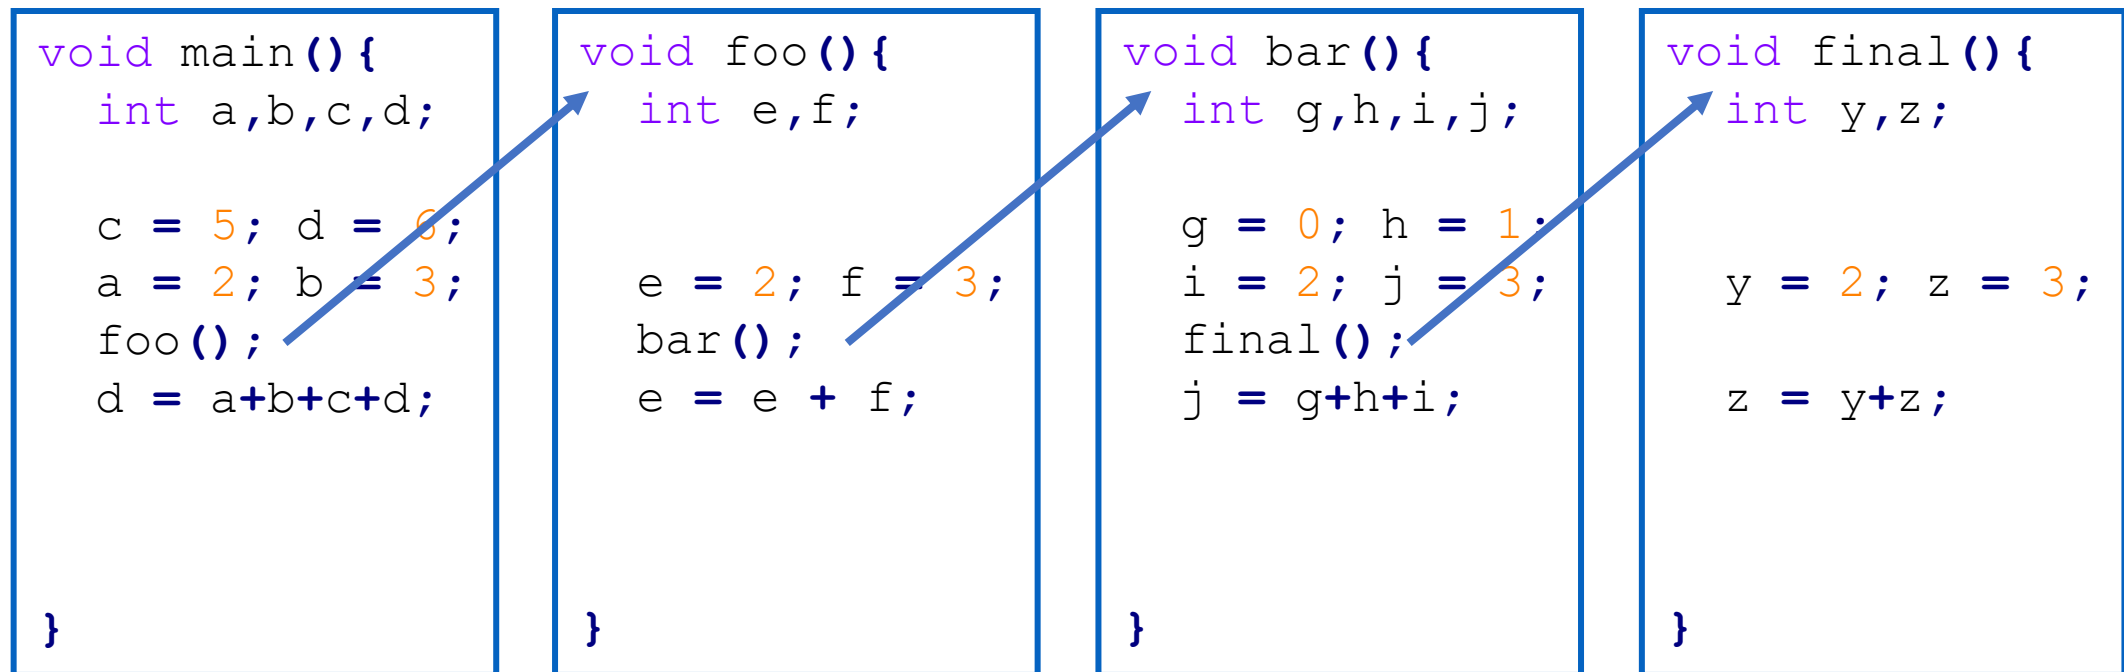
Agenda

- Branching far away
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- **Caller/callee example**

Caller versus Callee

- Which is better??
- Let's look at some examples...
- Simplifying assumptions:
 - A function can be invoked by many different call sites in different functions.
 - Assume no inter-procedural analysis (hard problem)
 - A function has no knowledge about which registers are used in either its caller or callee
 - Assume `main()` is not invoked by another function
- Implication
 - Any register allocation optimization is done using function local information

Caller-saved vs. callee saved – Multiple function case



Note: assume main does not have to save any callee registers

Caller-saved vs. callee saved – Multiple function case

- Questions:

1. How many registers need to be saved/restored if we use a **caller-save** convention?
2. How many registers need to be saved/restored if we use a **callee-save** convention?
3. How many registers need to be saved/restored if we use a mix of **caller-save** and **callee-save**?

Question 1: Caller-save

```
void main() {  
    int a,b,c,d;  
    c = 5; d = 6;  
    a = 2; b = 3;  
    [4 STUR]  
    foo();  
    [4 LDUR]  
    d = a+b+c+d;  
}
```

```
void foo() {  
    int e,f;  
  
    e = 2; f = 3;  
    [2 STUR]  
    bar();  
    [2 LDUR]  
    e = e + f;  
}
```

```
void bar() {  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    [3 STUR]  
    final();  
    [3 LDUR]  
    j = g+h+i;  
}
```

```
void final() {  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
}
```

Total: 9 STUR / 9 LDUR

Question 2: Callee-save

Poll: How many ld/st pairs are needed?

```
void main() {  
    int a,b,c,d;  
  
    c = 5; d = 6;  
    a = 2; b = 3;  
    foo();  
    d = a+b+c+d;  
}
```

```
void foo() {  
    [2 STUR]  
    int e,f;  
  
    e = 2; f = 3;  
    bar();  
    e = e + f;  
  
    [2 LDUR]  
}
```

```
void bar() {  
    [4 STUR]  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    final();  
    j = g+h+i;  
  
    [4 LDUR]  
}
```

```
void final() {  
    [2 STUR]  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
  
    [2 LDUR]  
}
```

Total: 8 STUR / 8 LDUR

Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
 - Use both!
 - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
 - Now each function can optimize for each situation to reduce saving/restoring
 - Not discussed further in this class

LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you want your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
 - X30 is the **link register** – used to hold return address
 - X28 is **stack pointer** – holds address of top of stack
 - X19-X27 are **callee-saved** – function must save these before writing to them
 - X0-15 are **caller-saved** –function must save live values before call
 - X0-X7 used for **arguments** (memory used if more space is needed)
 - X0 used for **return value**

Next Time

- Finish Up Function Calls
- Talks about linking – the final puzzle piece of software