

EECS 370 - Lecture 7

Linking



Announcements

- P1
 - Project 1 s + m due Thu
 - Instructor assembler available on the AG
- HW 1
 - Due Monday (9/22)
- Lab 4 meets Fr/M
 - Pre-Lab 4 quiz due Thursday
- Get exam conflicts sent to us **ASAP**
 - Forms listed on Ed

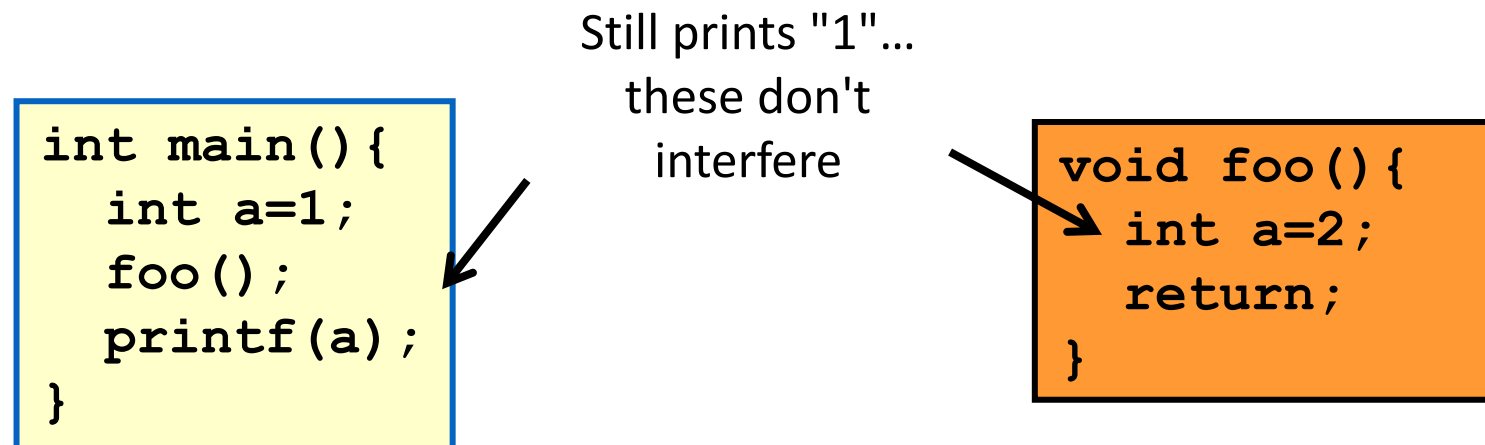
Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- **Lecture 7 : Translation software; libraries, memory layout**



Problem 3: Reusing registers

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
 - Even if different function have local variables with the same name, they are independent and guaranteed not to interfere with each other!



What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
 - Called functions will overwrite registers needed by calling functions

```
main: movz X0, #1  
      bl  foo  
      bl  printf
```

foo() overwrites
X0 if we don't
do something!!

```
foo: movz X0, #2  
     br  X30
```

- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee** saved)...

```
main: movz X0, #1
      bl  foo
      bl  printf
```

```
foo:  stur X0, [stack]
      movz X0, #2
      ldur X0, [stack]
      br  X30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller** saved)...

```
main: movz X0, #1
      stur X0, [stack]
      bl  foo
      ldur X0, [stack]
      bl  printf
```

```
foo:  movz X0, #2
      br  X30
```

Another example

Original C Code

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

No need to
save r2/r3.
Why?

Additions for Caller-save

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r4  
    restore r1  
    d = a+d;  
    return();  
}
```

Assume bar() will
overwrite registers
holding a,d

Additions for Callee-save

```
void foo() {  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r4  
    restore r3  
    restore r2  
    restore r1  
    return();  
}
```

bar() will save a,b, but
now foo() must save
main's variables

“caller-save” vs. “callee-save”

- Caller-save

- What if bar() doesn't use r1/r4?
- No harm done, but wasted work

```
void foo(){
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;
    save r1 to stack
    save r4 to stack
    bar();
    restore r1
    restore r4
    d = a+d;
    return();
}
```

- Callee-save

- What if main() doesn't use r1-r4?
- No harm done, but wasted work

```
void foo(){
    int a,b,c,d;
    save r1
    save r2
    save r3
    save r4
    a = 5; b = 6;
    c = a+1; d=c-1;
    bar();
    d = a+d;
    restore r1
    restore r2
    restore r3
    restore r4
    return();
}
```


Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- **Caller-saved**
 - Only needs saving if value is “live” across function call
 - **Live** = contains a useful value: Assign value before function call, use that value after the function call
 - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT
live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
 - Only needs saving at beginning of function and restoring at end of function
 - Only save/restore it if function overwrites the register

Only use r1-
r4

No need to
save other
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

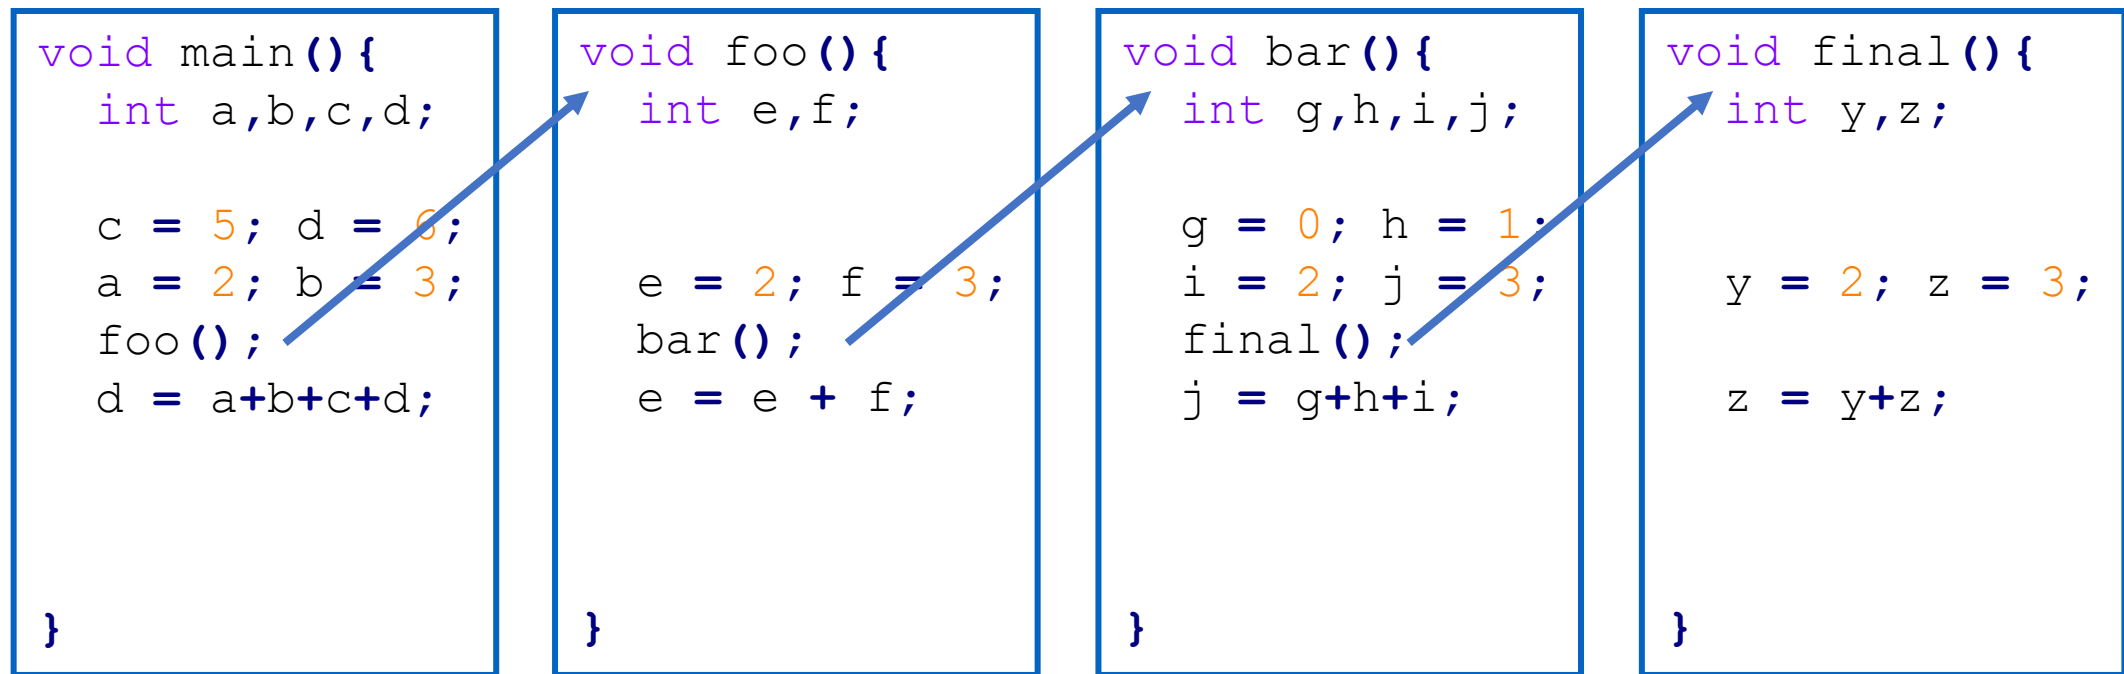
Agenda

- Branching far away
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- **Caller/callee example**

Caller versus Callee

- Which is better??
- Let's look at some examples...
- Simplifying assumptions:
 - A function can be invoked by many different call sites in different functions.
 - Assume no inter-procedural analysis (hard problem)
 - A function has no knowledge about which registers are used in either its caller or callee
 - Assume `main()` is not invoked by another function
- Implication
 - Any register allocation optimization is done using function local information

Caller-saved vs. callee saved – Multiple function case



Note: assume main does not have to save any callee registers

Caller-saved vs. callee saved – Multiple function case

- Questions:

1. How many registers need to be saved/restored if we use a **caller-save** convention?
2. How many registers need to be saved/restored if we use a **callee-save** convention?
3. How many registers need to be saved/restored if we use a mix of **caller-save** and **callee-save**?

Question 1: Caller-save

```
void main() {  
    int a,b,c,d;  
    c = 5; d = 6;  
    a = 2; b = 3;  
    [4 STUR]  
    foo();  
    [4 LDUR]  
    d = a+b+c+d;  
}
```

```
void foo() {  
    int e,f;  
  
    e = 2; f = 3;  
    [2 STUR]  
    bar();  
    [2 LDUR]  
    e = e + f;  
}
```

```
void bar() {  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    [3 STUR]  
    final();  
    [3 LDUR]  
    j = g+h+i;  
}
```

```
void final() {  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
}
```

Total: 9 STUR / 9 LDUR

Question 2: Callee-save

```
void main() {  
    int a,b,c,d;  
  
    c = 5; d = 6;  
    a = 2; b = 3;  
    foo();  
    d = a+b+c+d;  
  
}
```

```
void foo() {  
    [2 STUR]  
    int e,f;  
  
    e = 2; f = 3;  
    bar();  
    e = e + f;  
  
    [2 LDUR]  
}
```

```
void bar() {  
    [4 STUR]  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    final();  
    j = g+h+i;  
  
    [4 LDUR]  
}
```

```
void final() {  
    [2 STUR]  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
  
    [2 LDUR]  
}
```

Total: 8 STUR / 8 LDUR

Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
 - Use both!
 - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
 - Now each function can optimize for each situation to reduce saving/restoring
 - Not discussed further in this class

LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you want your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
 - X30 is the **link register** – used to hold return address
 - X28 is **stack pointer** – holds address of top of stack
 - X19-X27 are **callee-saved** – function must save these before writing to them
 - X0-15 are **caller-saved** –function must save live values before call
 - X0-X7 used for **arguments** (memory used if more space is needed)
 - X0 used for **return value**

Caller/Callee

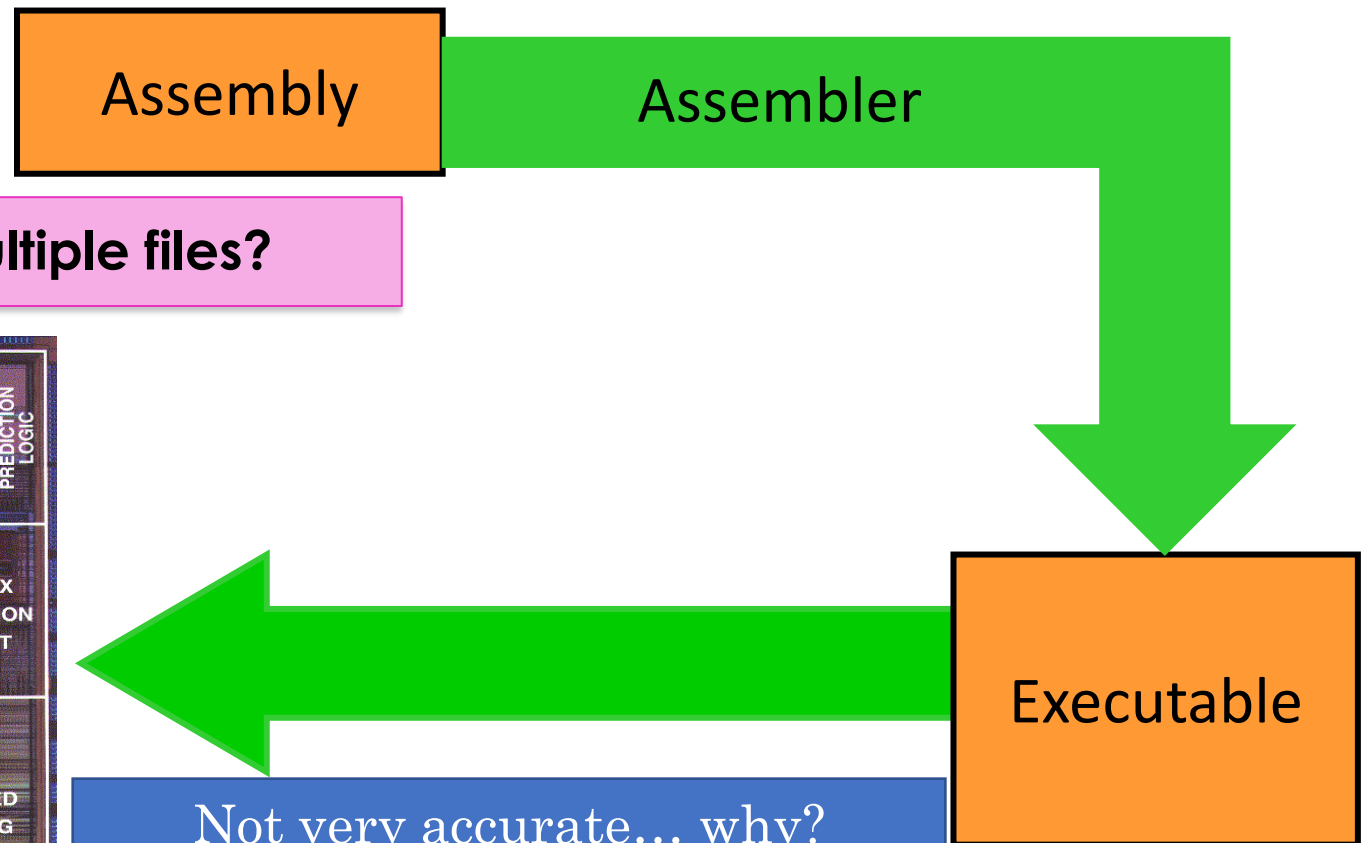
- Still not clicking?
- Don't worry, this is a tricky concept for students to get
- Check out supplemental video
 - <https://www.youtube.com/watch?v=SMH5uL3HiiU>
 - Don't worry about mixed caller/callee save case at the end
- Come to office hours to go over examples

Today we'll finish up software

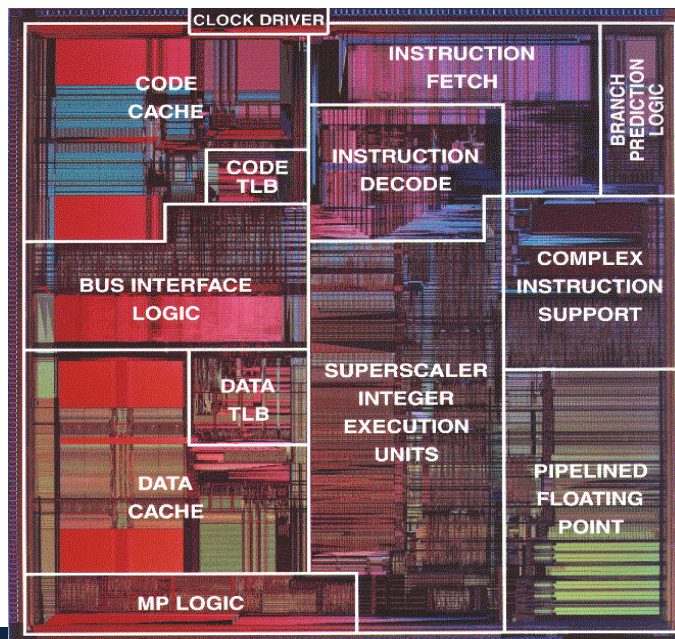
- Introduce linkers and loaders
 - Basic relationship of compiler, assembler, linker and loader.
 - Object files
 - Symbol tables and relocation tables

Source Code to Execution

- In project 1a, our view is this:



Why do we write programs in multiple files?




Not very accurate... why?
Because in reality, we have
multiple files

Multi-file programs

- In practice, programs are made from thousands or millions of lines of code
 - Use pre-existing libraries like stdlib
- If we change one line, do we need to recompile the whole thing?
 - No! If we compile each file into a separate **object file**, then we only need to recompile that one file and **link** it to the other, unchanged object files

Source Code to Execution

What do object files look like?



What do object files look like?

```
extern int X;  
extern void foo();  
int Y;  
  
void main() {  
    Y = X + 1;  
    foo();  
}
```

"extern" means
defined in another
file

```
extern int Y;  
int X;  
  
void foo() {  
    Y *= 2;  
}
```

Compile →

```
.main:  
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

Compile →

```
.foo:  
LDUR    X1, [XZR, Y]  
LSL     X9, X1, #1  
STUR    X9, [XZR, Y]  
BR      X30
```

Uh-oh!
Don't know
address of X, Y,
or foo!

Linking

```
.main:  
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

```
.foo:  
LDUR    X1, [XZR, Y]  
LSL     X9, X1, #1  
STUR    X9, [XZR, Y]  
BR      X30
```

Assemble

???

Assemble

???

What needs to go
in this intermediate
"object file"?

LINK

LINK

NOTE: this will
actually be in
machine code, not
assembly

```
LDUR    X1, [XZR, #40]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, #36]  
BL      #2  
HALT  
LDUR    X1, [XZR, #36]  
LSL     X9, X1, #1  
STUR    X9, [XZR, #36]  
BR      X30  
// Addr #36 starts here
```

Linking

.main:

```
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

Assemble

???

LINK

```
LDUR    X1, [XZR, #40]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, #36]  
BL      #2  
HALT  
LDUR    X1, [XZR, #36]  
LSL     X9, X1, #1  
STUR    X9, [XZR, #36]  
BR      X30
```


We need:

- the assembled machine code:
- list of instructions that need to be updated once addresses are resolved
- list of symbols for cross-ref

What do object files look like?

- Since we can't make executable, we make an object file
- Basically, includes the machine code that will go in the executable
 - Plus extra information on what we need to modify once we stitch all the other object files together
- Looks like this ->

Object code format



We won't discuss "Debug" much. Gets included when you compile with "-g" in gcc

Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    int Y = G + 1;
    B();
}
```

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
Reloc table	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B



Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    int Y = G + 1;
    B();
}
```

Header:
keeps track of
size of each
section

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
Reloc table	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    int Y =
    B();
}
```

Text:
machine code

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
Reloc table	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

Simplifying Assumption for EECS370

All globals and static locals (initialized or not) go in the data segment

Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    int Y = G + 1;
    B();
}
```

Data:
initialized globals
and static locals

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

Header	Name	foo		
	Text size	0x0C //probably bigger		
	Data size	0x04 //probably bigger		
Text	Address	Instruction		
	0	LDUR X1, [XZR, G]		
	4	ADDI X9, X1, #1		
	8	BL B		
Data	0	X	3	
Symbol table	Label	Address		
	X	0		
	B	-		
	main	0		
	G	-		
Reloc table	Addr	Instruction type	Dependency	
	0	LDUR	G	
	8	BL	B	



Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    int Y = G + 1;
    B();
}
```

LDUR
ADDI
BL

Symbol table:
Lists all labels
visible outside this file
(i.e. function names
and global variables)

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
Reloc table	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    int Y = G + 1;
    B();
}
```

LDUR X1, [XZR, G]

Relocation Table:
list of instructions and data words that must be updated if things are moved in memory

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
Reloc table	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

Class Problem 1

Which symbols will be put in the symbol table? (i.e. which "things" should be visible to all files?)

```
file1.c
extern void bar(int);
extern char c[];
int a;
int foo (int x) {
    int b;
    a = c[3] + 1;
    bar(x);
    b = 27;
}
```

file 1 – symbol table

sym	loc
a	data
foo	text
c	-
bar	-

```
file2.c
extern int a;
char c[100];
void bar (int y) {
    char e[100];
    a = y;
    c[20] = e[7];
}
```

file 2 – symbol table

sym	loc
c	data
bar	text
a	-

Which lines / instructions are in the relocation table? (i.e. which "things" need to be updated after linking?)

Class Problem 2

file1.c

```
1 extern void bar(int);
2 extern char c[];
3 int a;
4 int foo (int x) {
5     int b;
6     a = c[3] + 1;
7     bar(x);
8     b = 27;
9 }
```

file 1 - relocation table

line	type	dep
6	ldur	c
6	stur	a
7	bl	bar

file2.c

```
1 extern int a;
2 char c[100];
3 void bar (int y) {
4     char e[100];
5     a = y;
6     c[20] = e[7];
7 }
```

file 2 - relocation table

line	type	dep
5	stur	a
6	stur	c

Note: in a real relocation table, the "line" would really be the address in "text" section of the instruction we need to update.

Linker

- Stitches independently created object files into a single executable file (i.e., a.out)
 - Step 1: Take text segment from each .o file and put them together.
 - Step 2: Take data segment from each .o file, put them together, and concatenate this onto end of text segments.
- What about libraries?
 - Libraries are just special object files.
 - You create new libraries by making lots of object files (for the components of the library) and combining them (see ar and ranlib on Unix machines).
- Step 3: Resolve cross-file references to labels
 - Make sure there are no undefined labels

Linker - Continued

- Determine the memory locations the code and data of each file will occupy
 - Each function could be assembled on its own
 - Thus, the relative placement of code/data is not known up to this point
 - **Must relocate absolute references to reflect placement by the linker**
 - PC-Relative Addressing (beq, bne): never relocate
 - Absolute Address (mov 27, #X): always relocate
 - External Reference (usually bl): always relocate
 - Data Reference (often movz/movk): always relocate
- Executable file contains no relocation info or symbol table
these just used by assembler/linker

Loader

- Executable file is sitting on the disk
- Puts the executable file code image into memory and asks the operating system to schedule it as a new process
 - Creates new address space for program large enough to hold text and data segments, along with a stack segment
 - Copies instructions and data from executable file into the new address space
 - Initializes registers (PC and SP most important)
- Take operating systems class (EECS 482) to learn more!

Summary

- Compiler converts a single source code file into a single assembly language file
- Assembler handles directives (.fill), converts what it can to machine language, and creates a checklist for the linker (relocation table). This changes each .s file into a .o file
- Assembler does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses
- Linker enables separate compilation: Thus unchanged files, including libraries need not be recompiled.
- Linker resolves remaining addresses.
- Loader loads executable into memory and begins execution

Floating Point Arithmetic




Why floating point

- Have to represent real numbers somehow
- Rational numbers
 - Ok, but can be cumbersome to work with
- Fixed point
 - Do everything in thousandths (or millionths, etc.)
 - Not always easy to pick the right units
 - Different scaling factors for different stages of computation
- **Scientific notation: this is good!**
 - Exponential notation allows HUGE dynamic range
 - Constant (approximately) relative precision across the whole range

IEEE Floating point format (single precision)

- Sign bit: (0 is positive, 1 is negative)
- Significand: (also called the *mantissa*; stores the 23 most significant bits after the decimal point)
- Exponent: used biased base 127 encoding
 - Add 127 to the value of the exponent to encode:
 - -127 → 00000000 1 → 10000000
 - -126 → 00000001 2 → 10000001
 - ...
 - 0 → 01111111 128 → 11111111
- How do you represent zero ? Special convention:
 - Exponent: -127 (all zeroes), Significand 0 (all zeroes), Sign + or -

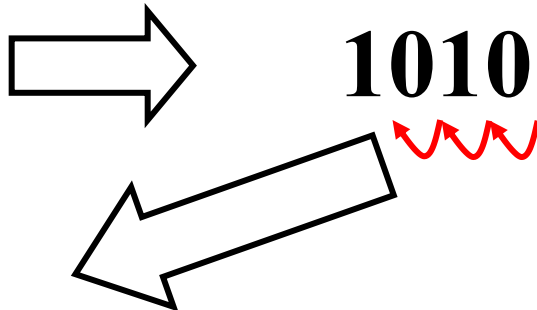


Floating Point Representation

$$10.625_{10} \Rightarrow 1010.101_2$$

- Step 1: convert from decimal to binary
 - 1st bit after "binary" point represents 0.5 (i.e. 2^{-1})
 - 2nd bit represents 0.25 (i.e. 2^{-2})
 - etc.

Floating Point Representation

$$10.625_{10} \quad \Rightarrow \quad 1010.101_2$$


$$1.010101 \times 2^3$$


- Step 2: normalize number by shifting binary point until you get $1.XXX \times 2^Y$

Floating Point Representation

$$10.625_{10} \longrightarrow 1010.101_2$$

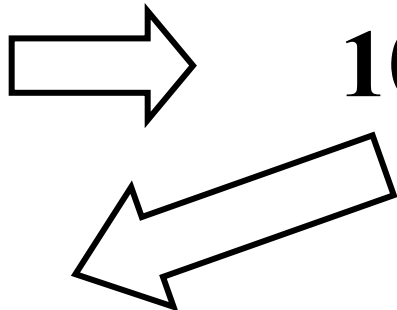
$$1.010101 \times 2^3$$

This must be a 1!
So don't store it.



- Step 3: store relevant numbers in proper location (ignoring initial 1 of significand)

Floating Point Representation

$$10.625_{10} \longrightarrow 1010.101_2$$


$$1.010101 \times 2^3$$

**This must be a 1!
So don't store it.**



$$10.625_{10} = 0 \ 10000010 \ 010101000000000000000000$$

Next Time

- Wrap up Floating Point
- And... hardware time!